


This repository has been archived by the owner. It is now read-only.

Branch: master ▾

Find file

Copy path



STAT545-UBC-original-website / automation01_slides / slides.md

 **jennybc** major reorganization of the automation content
c2cab17 on 7 Nov 2014

1 contributor


RawBlameHistory

649 lines (438 sloc)14.1 KB



title	author	date
Automating Data-analysis Pipelines	Shaun Jackman	2014-11-03

Automating Data-analysis Pipelines

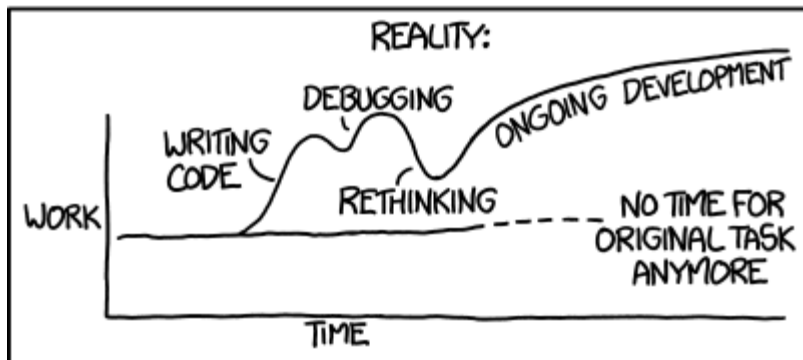
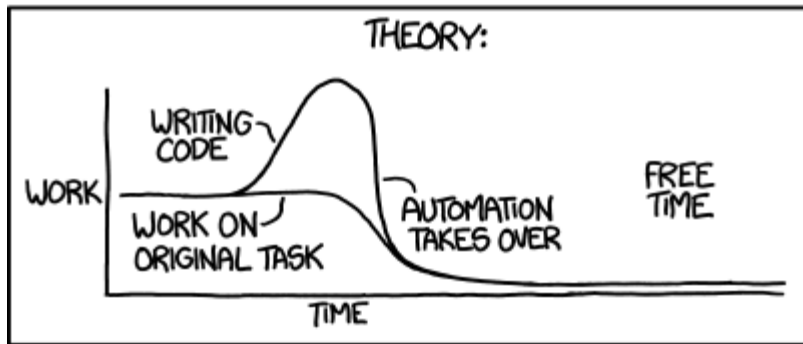
| UBC [STAT 545A](#)/STAT 547M | 2014-11-03 | [Shaun Jackman @sjackman](#) | [Jenny Bryan @JennyBryan](#) | 

Pipelines Automation Dependencies

Automation

- using [R](#)
- using the [shell](#) and [Rscript](#)
- using [make](#)

"I SPEND A LOT OF TIME ON THIS TASK.
I SHOULD WRITE A PROGRAM AUTOMATING IT!"



'Automating' comes from the roots 'auto-' meaning 'self-', and 'mating', meaning 'screwing'.

Pipelines

A pipeline

breaks up a monolithic make-all-the-things script into discrete, manageable chunks.

Each stage of the pipeline

| ... defines its input and its outputs. | ... does *not* modify its inputs, so it is *idempotent*.

| Rerunning a stage of the pipeline | produces the same results as the previous run.

Advantage #1

| When you modify one stage of the pipeline, | you don't have to rerun the entire pipeline.

You only rerun the downstream, dependent stages.

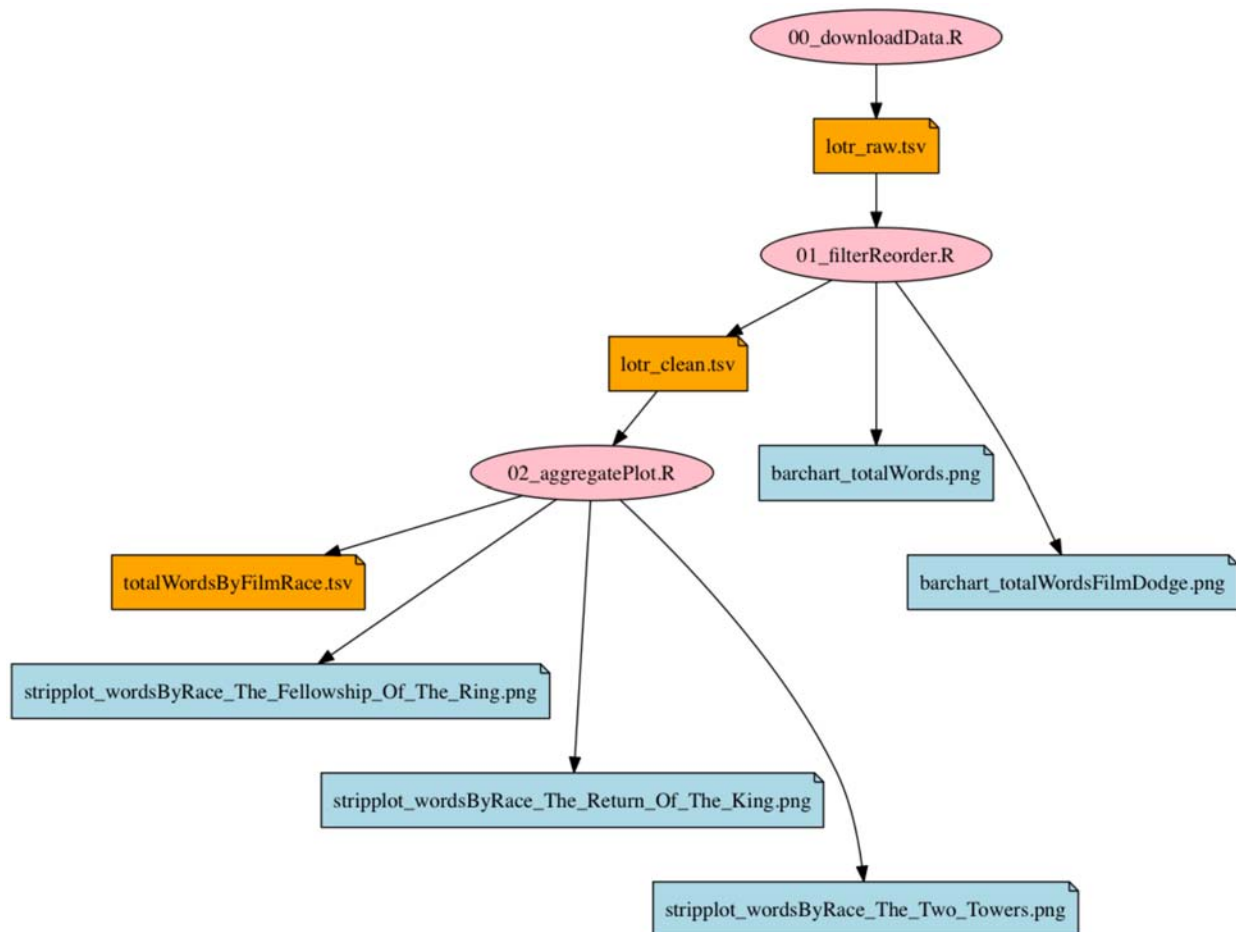
Advantage #2

Divide up work amongst a group by assigning to each person stages of the pipeline design.

Advantage #3

| You can draw pretty pictures of your pipeline, | because a pipeline is a graph.

|



| *01_justR*

Automation

Automate a pipeline

| ... to reproduce previous results. | ... to recreate results deleted by fat fingers. | ... to rerun the pipeline with updated software. | ... to run the same pipeline on a new data set.

Tools

- [R](#)
- the [shell](#) and [Rscript](#)
- [make](#)

R

```
#!/usr/bin/env Rscript
source("00_downloadData.R")
source("01_filterReorder.R")
source("02_aggregatePlot.R")
```

...

- Shows in what order to run the scripts.
- You can resume the pipeline from the middle.

Shell and Rscript

```
#!/bin/sh
set -eux
Rscript 00_downloadData.R
Rscript 01_filterReorder.R
Rscript 02_aggregatePlot.R
```

...

Allows you to easily run your pipeline from the shell.

...

Option	Effect
set -e	Stop at the first error
set -u	Undefined variables are an error
set -x	Print each command as it is run

Mix R scripts with other tools

```
#!/bin/sh
set -eux
curl -L http://bit.ly/lotr_raw-tsv >lotr_raw.tsv
```

```
Rscript 01_filterReorder.R  
Rscript 02_aggregatePlot.R
```

...

R is a good tool, but not always the best tool for the job.

Not sacrilege, but the principal tenet of a polyglot.

Makefile

```
#!/usr/bin/make -f  
  
lotr_raw.tsv:  
    curl -L http://bit.ly/lotr_raw-tsv >lotr_raw.tsv  
  
lotr_clean.tsv: 01_filterReorder.R lotr_raw.tsv  
    Rscript 01_filterReorder.R  
  
totalWordsByFilmRace.tsv: 02_aggregatePlot.R lotr_clean.tsv  
    Rscript 02_aggregatePlot.R
```

...

| A Makefile gives both the commands | *and their dependencies*.

Make is beautiful

| Tell Make how to create one type of file from another | and which files you want to create.

...

| Make looks at which files you have | and figures out how to create the files that you want.

Dependency graph

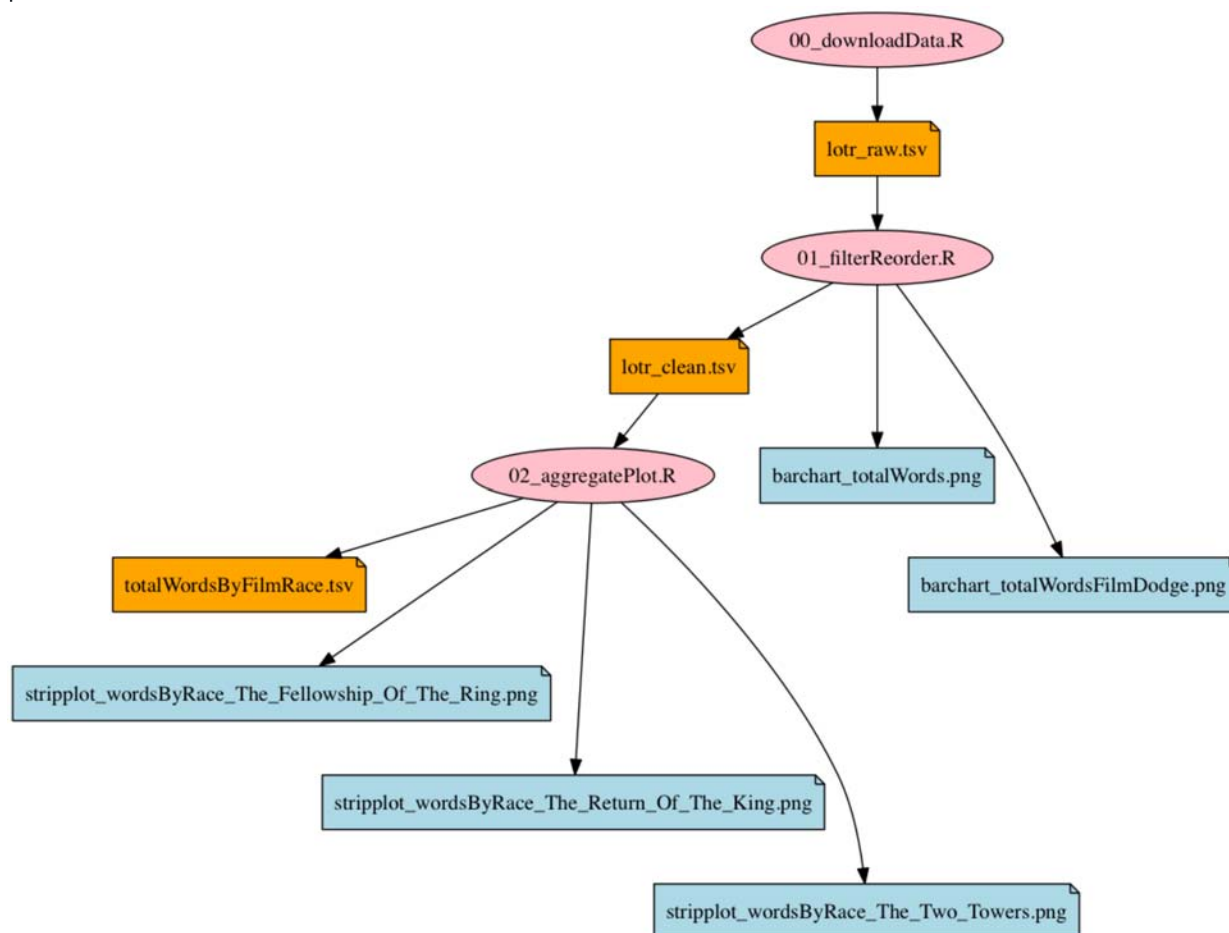
A pipeline is a graph

Scripts and data files are vertices of the graph.

Dependencies between stages are edges of the graph.

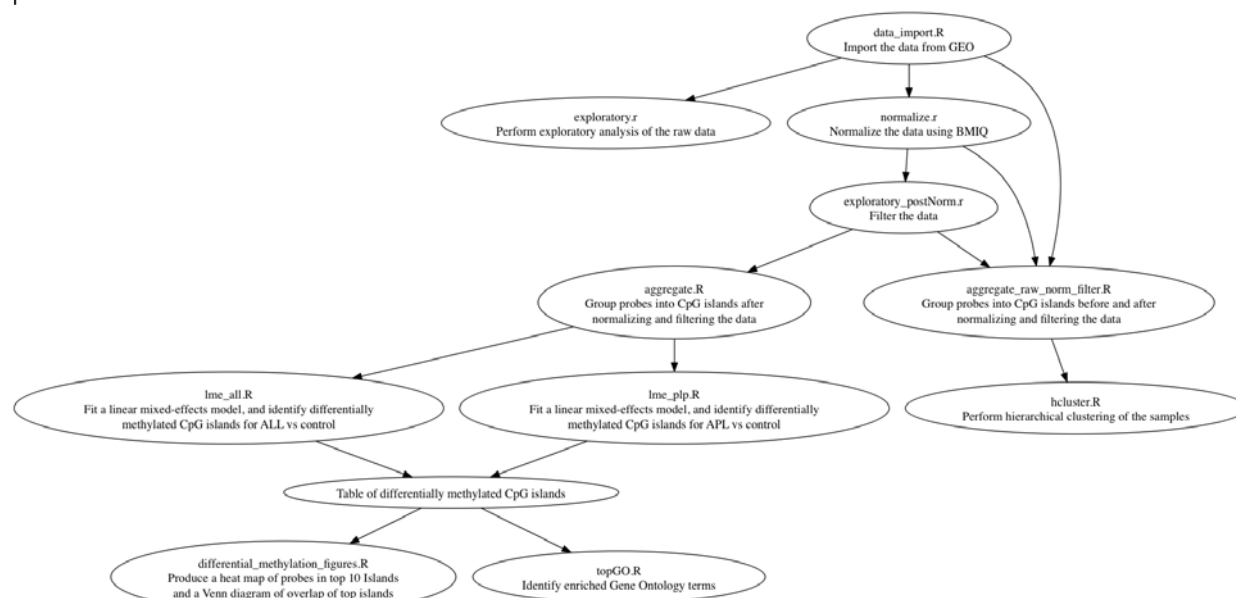
Both scripts and data files are shown.

|



| 01_justR

- Only dependencies between scripts are shown.
- Data files are not shown.
- Run the scripts in *topographical* order.



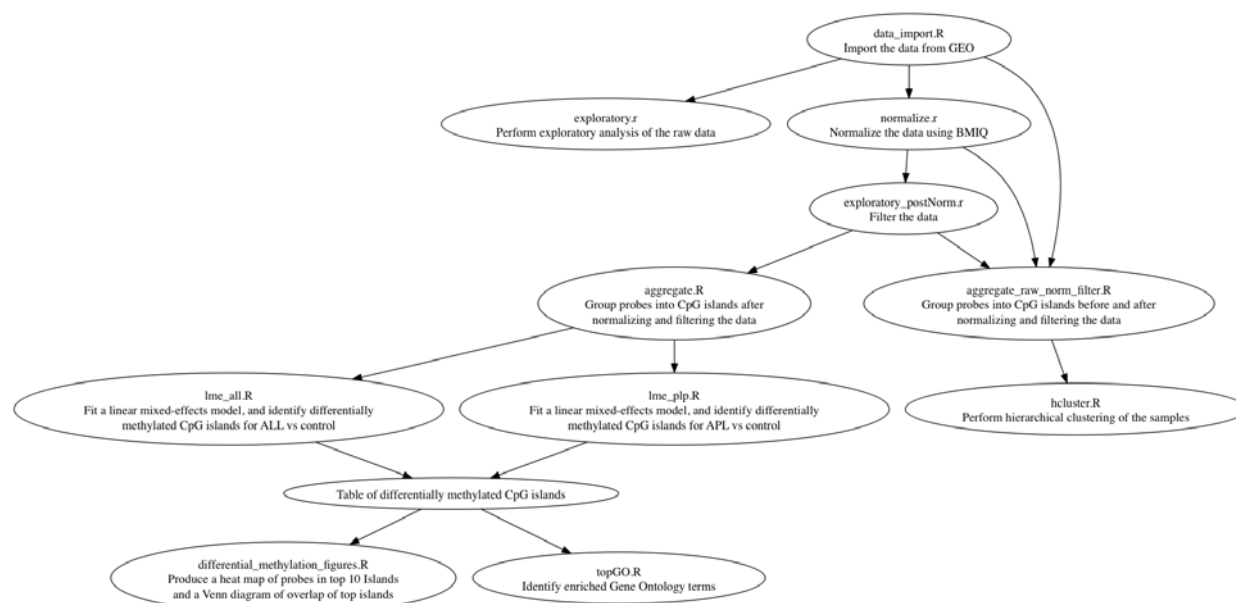
| STAT 540 Differential Methylation in Leukemia

Order of dependencies

A shell script gives *one* order in which you can successfully run the pipeline.

...

Unless the pipeline is completely linear, there are likely other such orders.



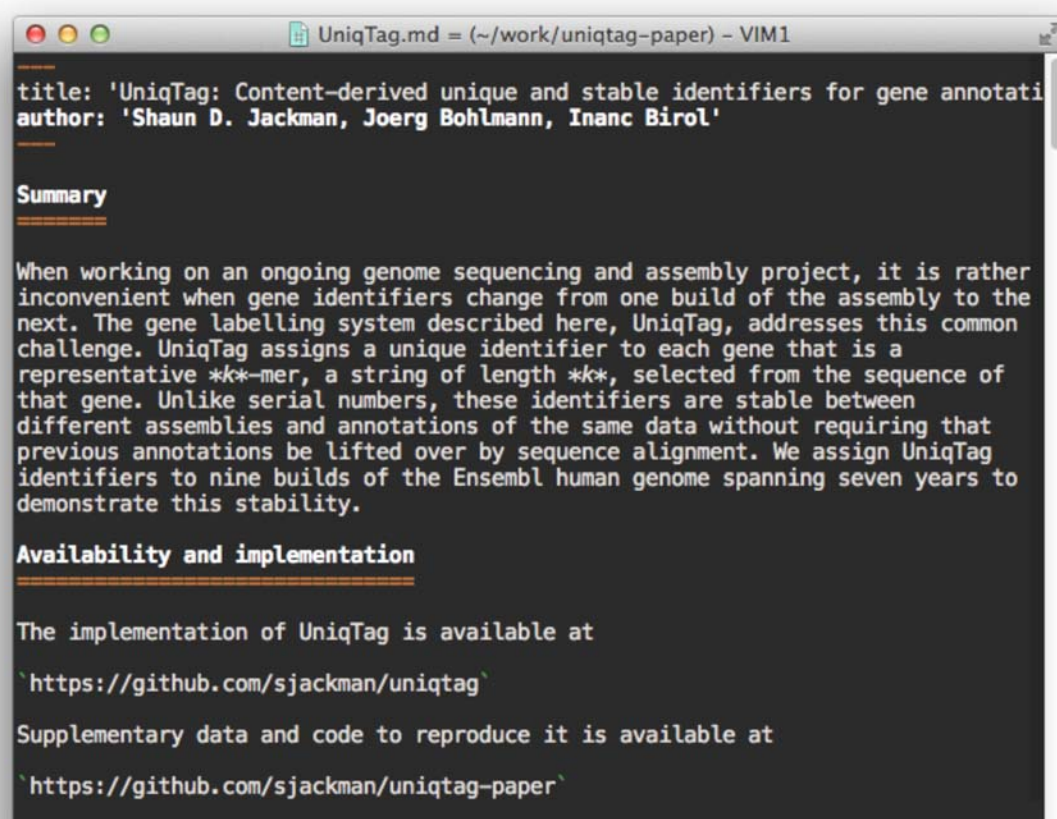
A different order of commands may be more convenient, but without information of the dependencies, you're stuck with the given order.

A reproducible manuscript

One Makefile

- Downloads the data
- Runs the command-line programs
- Performs the statistical analyses using [R](#)
- and Generates the TSV tables
- Renders the figures using [ggplot2](#)
- Renders the supplementary material using [RMarkdown](#)
- Renders the manuscript using [Pandoc](#)

Turns this



Into this

UniqTag: Content-derived unique and stable identifiers for gene annotation

Shaun Jackman^{1,2,*}, Joerg Bohlmann^{3,4} and İnanç Birol^{1,5 *}

¹Genome Sciences Centre, British Columbia Cancer Agency, Vancouver, BC, Canada

²Graduate Program in Bioinformatics, University of British Columbia, Vancouver, BC, Canada

³Michael Smith Laboratories, University of British Columbia, Vancouver, BC, Canada

⁴Department of Forest Science, University of British Columbia, Vancouver, BC, Canada

⁵Department of Medical Genetics, University of British Columbia, Vancouver, BC, Canada

Received on XXXXX; revised on XXXXX; accepted on XXXXX

Associate Editor: XXXXXXXX

ABSTRACT

Summary

When working on an ongoing genome sequencing and assembly project, it is rather inconvenient when gene identifiers change from one build of the assembly to the next. The gene labelling system described here, UniqTag, addresses this common challenge. UniqTag assigns a unique identifier to each gene that is a representative k -mer, a string of length k , selected from the sequence of that gene. Unlike serial numbers, these identifiers are stable between different assemblies and annotations of the same data without requiring that previous annotations be lifted over by sequence alignment. We assign UniqTag identifiers to nine builds of the Ensembl human genome spanning seven years to demonstrate this stability.

Availability and implementation

The implementation of UniqTag is available at

<https://github.com/sjackman/uniqtag>

Supplementary data and code to reproduce it is available at

<https://github.com/sjackman/uniqtag-paper>

Contact

Shaun Jackman <sjackman@bcgsc.ca>

İnanç Birol <ibirol@bcgsc.ca>

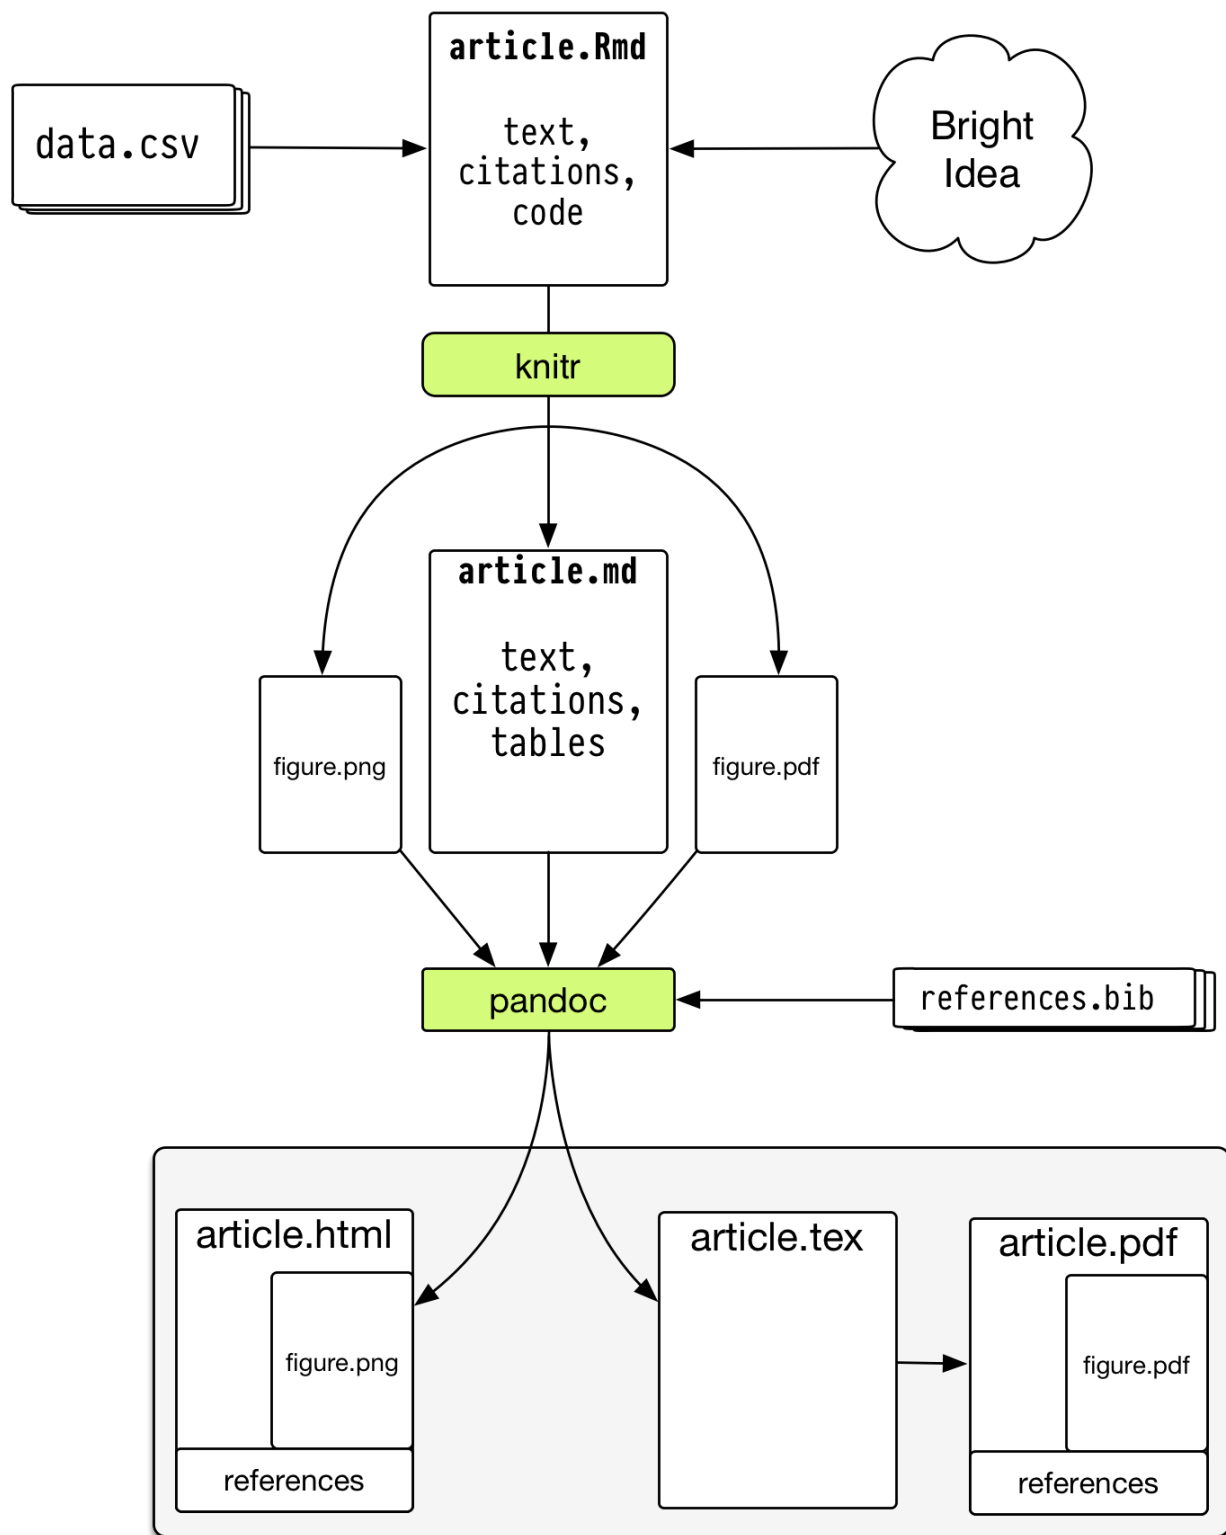
as SHA (Secure Hash Algorithm) (Dang, 2012) derives a message digest from the sequence, such that two sequences with the same content will have the same message digest, and two sequences that differ will have different message digests. If a cryptographic hash were used to identify a gene, the same gene in two assemblies with identical content would be assigned identical identifiers. Yet, by design a slight change in the sequence, such as a single-character substitution, would result in a completely different digest.

Locality-sensitive hashing in contrast aims to assign items that are similar to the same hash value. A hash function that assigns an identical identifier to a sequence after a modification of that sequence is desirable for labelling the genes of an ongoing genome annotation project. One such locality-sensitive hash function, MinHash, was employed to identify web pages with similar content (Broder, 1997) by selecting a small representative set of words from a web page.

UniqTag is inspired by MinHash. It selects a single representative k -mer from a sequence to assign a stable identifier to a gene. These identifiers are intended for systematic labelling of genes rather than assigning biological gene names, as the latter are typically based on biological function or homology to orthologous genes.

Workflow

Plain Text, Papers, Pandoc by Kieran Healy



Markdown for the manuscript

Markdown is a plain-text typesetting language

A header
=====

A list:

- + This text is **italic**
- + This text is ****bold****

A header

A list:

- This text is *italic*
- This text is **bold**

RMarkdown

- [RMarkdown](#) interleaves prose with [R](#) code
 - to aggregate and summarize the data
 - to generate tables
 - to render figures using [ggplot2](#)
- [RMarkdown](#) is ideal for supplementary material

RMarkdown example

```
The Sum of 1 + 1
=====
```

The sum of 1 + 1 is calculated as follows.

```
` `{r}
1 + 1
` `
```

```
![*Fig. 1*: A graphical view of 1 + 1](figure.png)
```

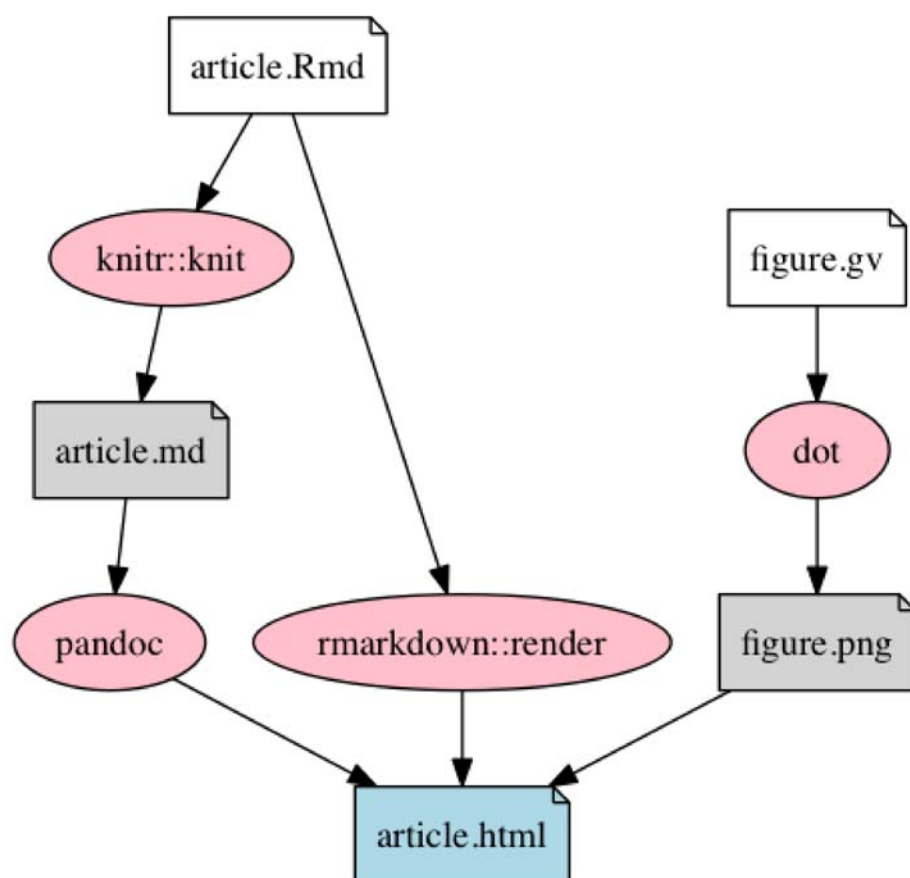
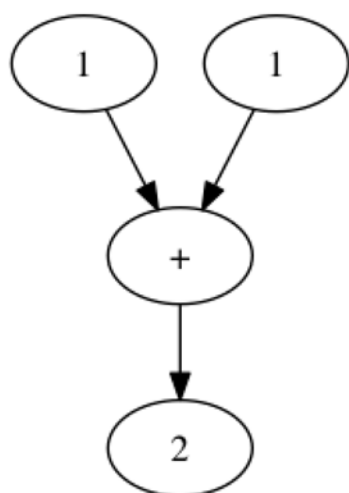
[article.Rmd](#)

The Sum of 1 + 1

The sum of 1 + 1 is calculated as follows.

```
1 + 1
```

```
## [1] 2
```



Dependencies of [article/Makefile](#)

Render HTML

```
%.md: %.Rmd
    Rscript -e 'knitr::knit("$<", "$@")'
```

```
%.html: %.md
    pandoc -s -o $@ $<
```

```
%.html: %.Rmd
    Rscript -e 'rmarkdown::render("$<")'
```

```
article.html: figure.png

%.png: %.gv
    dot -Tpng $< >$@

make article.html

dot -Tpng figure.gv >figure.png
Rscript -e 'rmarkdown::render("article.Rmd")'
```

Knit Markdown

```
%.md: %.Rmd
    Rscript -e 'knitr::knit("$<", "$@")'

%.html: %.md
    pandoc -s -o $@ $<

%.html: %.Rmd
    Rscript -e 'rmarkdown::render("$<")'

article.html: figure.png

%.png: %.gv
    dot -Tpng $< >$@

make article.md article.html

Rscript -e 'knitr::knit("article.Rmd", "article.md")'
dot -Tpng figure.gv >figure.png
pandoc -s -o article.html article.md
```

Pandoc

| [Pandoc](#) renders attractive documents and slides | from plain-text typesetting formats

It converts between every format known (just about)

- Markdown
- HTML
- LaTeX

- PDF
- ODT and docx (yes, really)

Evolving a Makefile

```
#!/bin/sh
set -eux
dot -Tpng -o figure.png figure.gv
Rscript -e 'knitr::knit("article.Rmd")'
pandoc -s -o article.html article.md
```

Shell script

```
all:
    dot -Tpng -o figure.png figure.gv
    Rscript -e 'knitr::knit("article.Rmd")'
    pandoc -s -o article.html article.md
```

First Makefile

```
all: article.html

article.html:
    dot -Tpng -o figure.png figure.gv
    Rscript -e 'knitr::knit("article.Rmd")'
    pandoc -s -o article.html article.md
```

Add a rule to build article.html

```
all: article.html

article.html: article.Rmd
    dot -Tpng -o figure.png figure.gv
    Rscript -e 'knitr::knit("article.Rmd")'
    pandoc -s -o article.html article.md
```

article.html depends on article.Rmd

```
all: article.html

figure.png: figure.gv
    dot -Tpng -o figure.png figure.gv

article.md: article.Rmd
    Rscript -e 'knitr::knit("article.Rmd")'

article.html: article.md figure.png
    pandoc -s -o article.html article.md
```

Split one rule into three

```
all: article.html

figure.png: figure.gv
    dot -Tpng -o $@ $<

article.md: article.Rmd
    Rscript -e 'knitr::knit("$<", "$@")'

article.html: article.md figure.png
    pandoc -s -o $@ $<
```

Use the variables `$<` and `$@` for the input and output file

```
all: article.html

%.png: %.gv
    dot -Tpng -o $@ $<

%.md: %.Rmd
    Rscript -e 'knitr::knit("$<", "$@")'

article.html: article.md figure.png
    pandoc -s -o $@ $<
```

Use pattern rules. The `%` matches any string

```
all: article.html

%.png: %.gv
```

```

dot -Tpng -o $@ $<

%.md: %.Rmd
    Rscript -e 'knitr::knit("$<", "$@")'

%.html: %.md
    pandoc -s -o $@ $<

article.html: figure.png

```

article.html also depends on figure.png

```

all: article.html

clean:
    rm -f article.md article.html figure.png

%.png: %.gv
    dot -Tpng -o $@ $<

%.md: %.Rmd
    Rscript -e 'knitr::knit("$<", "$@")'

%.html: %.md
    pandoc -s -o $@ $<

article.html: figure.png

```

Add the target named clean

```

all: article.html

clean:
    rm -f article.md article.html figure.png

.PHONY: all clean
.DELETE_ON_ERROR:
.SECONDARY:

%.png: %.gv
    dot -Tpng -o $@ $<

%.md: %.Rmd
    Rscript -e 'knitr::knit("$<", "$@")'

%.html: %.md
    pandoc -s -o $@ $<

```



```
article.html: figure.png
```

Add `.PHONY`, `.DELETE_ON_ERROR` and `.SECONDARY`

```
all: article.html
```

```
clean:
```

```
    rm -f article.md article.html figure.png
```

```
.PHONY: all clean
```

```
.DELETE_ON_ERROR:
```

```
.SECONDARY:
```

```
# Render a GraphViz file
```

```
%.png: %.gv
```

```
    dot -Tpng -o $@ $<
```

```
# Knit a RMarkdown document
```

```
%.md: %.Rmd
```

```
    Rscript -e 'knitr::knit("$<", "$@")'
```

```
# Render a Markdown document to HTML
```

```
%.html: %.md
```

```
    pandoc -s -o $@ $<
```

```
# Dependencies on figures
```

```
article.html: figure.png
```

fin

Links

| [STAT 545A](#) | [xkcd automation](#) | [R](#) | [Rscript](#) | [shell](#) | [make](#) | [Markdown](#) | [RMarkdown](#) | [Pandoc](#) | [ggplot2](#) | [Plain Text, Papers, Pandoc](#) | [STAT 540 Differential Methylation in Leukemia](#)

[@sjackman](#) | github.com/sjackman | sjackman.ca

Shaun Jackman

| [Genome Sciences Centre](#), BC Cancer Agency | Vancouver, Canada | [@sjackman](#) | github.com/sjackman | sjackman.ca

